

ClaimReady: Small Claims, Filed Right

For the invoice too small for a lawyer and too large to write off.

ClaimReady helps NYC freelancers turn messy contracts, invoices, emails, and screenshots into a court-ready small-claims packet for unpaid service work.

The user. The concrete user is a New York City freelancer or solo service provider owed \$1,000–\$10,000 by a NY-registered LLC or corporation: a designer, developer, photographer, writer, tutor, consultant, or marketer with a contract, invoice, emails, screenshots, or notes showing completed work and nonpayment. This is not a niche edge case. Upwork estimates that 64 million Americans freelanced in 2023, or 38% of the U.S. workforce. A New York freelancer survey found that 62% experienced nonpayment; among those who lost income, 51% lost more than \$1,000 and 22% lost more than \$5,000.

The problem. Today, the user sends awkward follow-up emails, copies a generic demand-letter template, gives up, or tries to file in small claims alone. Hiring counsel rarely makes sense: Clio reports the average New York lawyer hourly rate at \$426, which can consume the economics of a \$2,000–\$5,000 claim. Small claims is designed for this gap, but filing still requires the user to identify the correct legal defendant, find a service address, choose venue, organize exhibits, calculate damages, and create forms the clerk can process. NYC Small Claims allows claims up to \$10,000 with only a \$15–\$20 filing fee, but the paperwork and confidence gap stop many people before they start.

The economics. For one user-month, assume one case packet: about 120k input tokens for intake, evidence, retrieval, and tool context, plus 15k output tokens for extracted facts, agent outputs, and final drafting. With Gemini 2.5 Flash pricing around \$0.30 per 1M input tokens and \$2.50 per 1M output tokens, model cost is about \$0.07 per case. Cloud Run, storage, PDF generation, logs, failed runs, payment fees, and light support add an estimated \$3–\$5 of variable cost; the model is not the cost driver—infrastructure and Stripe fees are.

Pricing ladder. ClaimReady charges \$29 per packet at launch and layers in revenue capture as volume justifies the operational lift. **Phase 1** is flat \$29 to validate willingness-to-pay and accumulate the first thousand cases of testimonials and defendant patterns; the asymmetry it accepts is real (a \$1,000 claim pays a 2.9% rate while a \$10,000 claim pays only 0.3%), but a single price keeps the funnel simple while the product earns trust. **Phase 2** adds **1–2% of the claim value at filing** on top of the \$29 base, charged upfront rather than as a contingency on recovery so there is no collection risk and no unauthorized-practice-of-law framing issue; the trigger is roughly 500 cases/month sustained, by which point Stripe metering and value-based capture are worth the engineering work. **Phase 3** is B2B distribution: ClaimReady’s engine white-labeled inside freelancer-facing platforms—Upwork, Fiverr, FreshBooks, QuickBooks Self-Employed—at a \$10–\$15 wholesale price, with the partner paying acquisition and capturing the freelancer-retention benefit, unlocked once consumer-channel CAC and packet-quality metrics are proven.

Item	Phase 1 (Launch)	Phase 2 (Scale)	Phase 3 (Distribution)
Price	\$29 / packet	\$29 + 1.5% of claim value	\$10–\$15 wholesale
Avg revenue / case	\$29	~ \$81 (avg claim \$3.5k)	~ \$12
Variable cost	\$3–\$5	\$3–\$5	\$3–\$5
Contribution / case	~ \$25 (~85%)	~ \$76 (~94%)	~ \$8 (~67%)
CAC payback	< 1.5× / case	~3× / case	Partner-paid CAC

CAC and payback. Blended CAC at launch is estimated at \$15–\$30: SEO on the “how do I sue an LLC in NY” long-tail with borough-specific landing pages, presence in r/freelance and freelancer Twitter, and partnerships with coworking spaces and small-business accountants. That gives contribution-margin payback of under 1.5× per case at the \$29 flat price and roughly 3× per case once Phase 2 take-rate pricing is live. Per-case payback matters more than LTV here—small claims is fundamentally episodic, and a business that pays back every transaction does not depend on retention assumptions to be solvent.

Break-even and margin risks. Fixed monthly burn at launch is roughly \$8–\$12k (Cloud Run minimums, Vertex AI baseline, part-time-equivalent founder time), so break-even sits near 400 cases/month at flat \$29 and roughly 150 cases/month once Phase 2 is live—the take-rate transition is what makes the business compound rather than just survive. The margin killers are concentrated and known: huge evidence sets that explode token cost beyond the \$0.07 baseline, refund and chargeback exposure if packet quality slips, Stripe fees of about 4% on a \$29 ticket, high-touch support when service-of-process bounces, and jurisdiction-expansion engineering when the wedge moves outside NYC.

Why these technical choices. The agent design mirrors the real filing workflow. The Next.js wizard collects only the facts the court needs, reducing user confusion. The Extractor reads messy contracts, invoices, emails, and screenshots so users do not have to translate evidence into legal structure. The DefendantResolver calls the NY Department of State dataset because suing the wrong entity or using the wrong service address can sink the case. The JurisdictionChecker uses deterministic tools and a small Chroma RAG corpus to validate the \$10,000 cap, statute-of-limitations fit, venue, and citations without expensive broad research. Structured Pydantic outputs make the PDF renderer reliable. Gemini 2.5 Flash is the right model because it supports multimodal, long-context evidence review while keeping per-case token cost far below the \$29 price.

Risks and mitigations. *Bar and UPL risk:* ClaimReady is positioned as form preparation and educational guidance, not legal advice; every output carries an explicit disclaimer, and complex cases are flagged for attorney review rather than handled. *Data dependency:* the NY DOS Active Corporations API is single-source for defendant resolution today, mitigated by response caching, manual-entry fallback with verification, and a per-state plug architecture so expansion states can use their own registries. *Model and quality risk:* structured Pydantic outputs and deterministic tools for jurisdiction and damages bound the LLM’s surface area, RAG citations make legal claims auditable, and a packet-review checkpoint keeps the human in the loop before filing. *Refund and chargeback risk:* a “court-ready or your money back” guarantee is feasible because failure modes are concrete and observable, and the \$29 ticket caps fraud exposure relative to the trust-signal benefit.

Go-to-market. Three channels at launch with a partnership path layered on top.

1. SEO on the “how do I sue an LLC in NY” long-tail, with borough-specific landing pages and anonymized case-study posts.
2. Community presence in r/freelance, r/personalfinance, freelancer Twitter, and the Freelancers Union mailing list.
3. Local partnerships with coworking spaces, small-business accountants, and freelance-friendly co-ops, where ClaimReady is offered at a member-benefit discount.
4. The Phase 3 path: B2B distribution inside freelancer-facing platforms where the partner pays acquisition and ClaimReady captures the engine economics—unlocked once consumer-channel CAC and packet-quality metrics are proven.

MVP and wedge. The product as scoped is an MVP: NYC small-claims, unpaid-invoice disputes against NY-registered entities. That narrowness is the whole point. Winning one workflow gives ClaimReady real users, real packet data, real defendant patterns, and a defensible filing engine, all of which generalize. From there, expansion follows two axes. First, claim types: security-deposit recovery, breach of consumer-services contract, unreturned property, and small-business vendor disputes share the same intake-extract-validate-render skeleton and only require new templates and rule modules. Second, jurisdictions: every U.S. small-claims court has its own cap, venue rules, forms, and entity registry, but the agent architecture, JurisdictionChecker, and DefendantResolver swap cleanly per state. NYC is the wedge because it has the highest density of freelancers, a generous \$10,000 cap, and a public DOS entity dataset; once the unit economics are proven there, the same engine extends to CA, TX, IL, and FL with incremental rule work rather than a rewrite.

Close. ClaimReady wins by being deliberately small at launch: not “AI lawyer for everything,” but a cheap, focused filing assistant for one painful, common, economically underserved workflow, designed from day one as the wedge for a broader civil-claims platform.

Sources: Upwork Freelance Forward 2023; NY freelancer nonpayment survey; NYC Small Claims filing instructions; Clio NY lawyer rates; Google Vertex AI pricing.